

@shashikant.ai · LAYA

Laya vs Jev: **all the links**

Every source from the reel, so you can check it yourself, plus how to try Laya free.

The story in 30 seconds

- **Jev** is a "System One" decision model by TypeSafe AI (founder Diogo Almeida, ex-OpenAI), launched 15 Sep 2026. It doesn't write text. It returns typed decisions with a confidence score.
- **Laya** is a free, open-source model by Nandakishor Mukkunoth (ConvAI Innovations, Kasaragod, Kerala). It came out a few days later.

His claim: he published the same core idea in a research paper in March 2025. This is a claim, not proven. TypeSafe has not replied publicly, and TypeSafe says it worked in stealth for two years.

Sources

Jev (TypeSafe)

Launch post ("After two years in stealth")

typesafe.ai/blog/introducing-system-one-models-and-jev

Models and pricing (\$0.042 per 1M input tokens, output free)

docs.typesafe.ai/models

Vercel: fastest-adopted model in AI Gateway history

vercel.com/blog/ai-gateway-jev-model-launch

Laya and the claim

Laya on GitHub (Apache 2.0)

github.com/NandhaKishorM/laya

His article: "I built ... a year ago"

laya.convaiinnovations.com

Paper 1: SalesRLAgent (submitted 30 Mar 2025)

arxiv.org/abs/2503.23303

Paper 2: Confidence-Aware Routing (submitted 23 Sep 2025)

arxiv.org/abs/2510.01237

The Hacker News debate

news.ycombinator.com/item?id=49765348

News coverage (Startup Fortune)

startupfortune.com

Who is better

Independent test on 741 real documents: Jev 0.919 vs Laya 0.780

github.com/NandhaKishorM/laya/issues/450

Laya's own benchmark table (incl. Banking77)

github.com/NandhaKishorM/laya/blob/main/BENCHMARKS.md

Quick comparison

	Jev	Laya
Price	Paid (\$0.042 / 1M input tokens)	Free (Apache 2.0)

Model	Closed, API only	Open code + weights
Runs on	TypeSafe cloud	Your laptop or PC (CPU or GPU)
Speed	236–276 ms per question (independent)	~6–7x faster (Laya's own test, T4 GPU)
Accuracy, independent test	0.919	0.780
70+ options (Banking77)	0.870 (72 labels)	0.425 (77 labels)

Want it free and local? Use Laya. Need the most accuracy? Use Jev.

Try Laya free (Python 3.10+)

```
python -m pip install laya
```

```
from laya import Router

router = Router()
state = "Hi, we were billed twice for March. Please refund the duplicate."
questions = {
    "department": {"type": "choice", "instructions": "Which department?",
                  "criteria": {"billing": "invoices, payments, refunds",
                              "technical": "bugs, outages"}}
}
result = router.predict(state, questions)
print(result["answers"]["department"]["choice"])  # billing
```

It runs on a normal CPU (slower, a few hundred ms per question) or a GPU. Test it on your own data before you trust it for anything important.

Checked on 30 Sep 2026. Follow [@shashikant.ai](#) for more AI tools, explained honestly.